

Statistical Certification of Autonomous Vehicle Fallback Triggers: Closed-Loop Evaluation Under Dataset Shift

Hassan Waqas Saleem, Rahat Zeeshan Saleem, and Hina Maryam

Abstract—A fallback trigger decides when an automated vehicle should abandon its task for a minimum-risk manoeuvre (MRM). In practice it is a hand-tuned threshold with no stated guarantee, and statistical certification of such thresholds has been assessed on predictions or static labels rather than on what the fallback then does. We ask how reliable a certified trigger is when judged through actual interventions, and what happens to it under dataset shift. A rollout depends only on the tick at which the fallback fires, so one reference run and one forced-firing rollout per decision tick evaluate every trigger, threshold, online method and latency exactly, on identical physics. On replayed Argoverse 2 and nuScenes logs with reactive traffic, a UN R157 manoeuvre and nuPlan harm labels, we compare a Learn-then-Test-certified geometric trigger with seven baselines. At a 5 % target, the certified trigger misses 3.7 % of harmful events and exceeds the target in 9.5 % of 200 resplits (limit 10 %); the same score tuned without a certificate exceeds it in 43.5 %. That validity costs 2.5 points of unnecessary stops (28.5 % versus 26.0 %), so we do not claim a better trade-off. Confidence and ensemble triggers need 68–71 % unnecessary stops for comparable safety. A threshold certified on Argoverse 2 misses 13.2 % on nuScenes, 5.0 % of nuScenes scenarios contain harm no trigger can prevent, and one decision tick (0.5 s) of latency already breaks the target. In an exploratory analysis, a few dozen labelled target scenes restore a valid certificate but at a large cost in stops. [Waymo replication to be inserted.]

Index Terms—Autonomous driving, conformal prediction, selective prediction, minimum-risk manoeuvre, closed-loop safety evaluation, trajectory prediction, risk calibration, distribution shift.

I. INTRODUCTION

AUTONOMOUS vehicles that lack confidence in their own situational understanding are expected to fall back on a minimum-risk manoeuvre: slow to a stop, pull to the shoulder, or otherwise hand control back to a human. The decision of *when* to invoke this behaviour is, in nearly every deployed and published system, a hand-tuned threshold on some scalar confidence or uncertainty score produced by a perception or prediction module. That practice has a specific and largely unexamined weakness: the very quantity the threshold

is applied to — a learned model’s uncertainty estimate — is exactly what has been shown to degrade under the kind of distribution shift a vehicle experiences when it enters a new city, a new weather condition, or simply a stretch of road unlike anything in its calibration set. A threshold tuned under one operating distribution carries no guarantee once that distribution has moved, and the failure is silent: the system continues to report a confidence number, and nothing in the number itself signals that the number is now wrong.

Conformal prediction offers a principled alternative: rather than a point confidence score, it produces a prediction region with a distribution-free, finite-sample coverage guarantee under exchangeability. A substantial and fast-growing literature has applied conformal prediction to trajectory forecasting and to planning under uncertainty in autonomous driving. Almost without exception, however, this literature is *open-loop*: it constructs a prediction region, reports its coverage and its size, and stops. Two things are invisible from that vantage point. First, a region’s effect on a downstream control decision — specifically, whether it is used to trigger a fallback — is not evaluated at all; open-loop coverage and closed-loop safety outcome are simply assumed to move together. Second, and more consequentially, *triggering the fallback is not free*. A vehicle that stops in a live lane because its uncertainty region grew large is not obviously safer than one that continues; an evaluation that reports only coverage, or only collision rate under the fallback, cannot see the trade-off between the two, and cannot see that a policy which triggers whenever it is asked to would trivially achieve a good miss rate while being useless.

This paper asks a narrower question than “which trigger is best”: *how reliable is a statistically certified fallback trigger when it is judged through the interventions it actually causes, and what happens to it when the deployment data differ from the calibration data?* We treat the trigger as a decision problem with a formal guarantee attached to the decision itself, and evaluate it where that guarantee has to hold: in closed loop, on real driving scenarios, with a threshold whose consequences change what the vehicle does next. Certifying a threshold on a risk score against missed events is not new; Luo *et al.* [8] do so with conformal prediction, and Conformal Decision Theory [12] calibrates decisions rather than predictions through a long-run *average* guarantee. We build on these and do not claim the certification step as ours. An MRM trigger nevertheless differs in kind from the settings they study. It is an irreversible, one-shot decision to

H. W. Saleem (corresponding author) is with the School of Electrical Engineering and Computer Science (SEECs), National University of Sciences and Technology (NUST), H-12, Islamabad 44000, Pakistan (e-mail: hsaleem.msee20seecs@seecs.edu.pk).

R. Z. Saleem is with the NUST Business School (NBS), National University of Sciences and Technology (NUST), H-12, Islamabad 44000, Pakistan (e-mail: rahat.mscm21nbs@student.nust.edu.pk).

H. Maryam is with the State Bank of Pakistan, I.I. Chundrigar Road, Karachi 74000, Pakistan (e-mail: hina.maryam@sbp.org.pk).

Manuscript in preparation.

abandon the task, its errors are paid in two currencies (missed interventions, and the harm of stopping itself), and a deployed vehicle needs its guarantee before the first mile.

We organise the study around four research questions:

- **RQ1.** Does a certificate keep its promise when measured on the closed-loop outcome, and what does keeping it cost in unnecessary stops and fallback-induced collisions?
- **RQ2.** Do open-loop failure-detection metrics choose the same trigger that closed-loop safety does?
- **RQ3.** What happens to a certificate under a real cross-dataset shift, and how much can be repaired with unlabelled or with a few labelled target scenarios?
- **RQ4.** How sensitive is the certificate to the timing of the intervention?

We make four contributions.

- 1) **An exact closed-loop evaluation protocol.** A rollout's outcome depends only on *when* the fallback fires. One reference run that records every trigger score, plus one rollout forced to fire at each decision tick, evaluates every trigger, threshold, online method and latency on identical physics, with the residual simulator drift measured and reported (Section III-E).
- 2) **What certification buys and what it costs, in closed loop.** We price each operating point in missed interventions, unnecessary stops and fallback-induced collisions. Certification is the only design that reliably keeps its target across resplits, at a measured cost in stops. It does not win the safety–conservatism trade-off against a tuned threshold, a hypothesis we pre-registered and report as refuted. Confidence and ensemble triggers, the deployed practice, need several times more unnecessary stops than a geometric score for comparable safety.
- 3) **The certificate under real dataset shift.** An Argoverse 2 certificate fails on nuScenes. The nuScenes data also expose a floor: harm present at the first decision tick, which no trigger can prevent and which can make a target unreachable. A density-ratio-weighted certificate that reports its effective sample size repairs part of the gap, and in an exploratory analysis the transfer is asymmetric.
- 4) **Operating-point findings for practitioners.** Latency of one decision tick breaks the target, route-following ego corridors make a trigger fire in every scenario, and open-loop AUROC ranks triggers only approximately as closed-loop stops do.

All hypotheses, thresholds and refutation rules were recorded in a version-controlled note before the data were collected. Refuted and unsupported hypotheses are reported as such, and analyses run after the pre-registered verdicts are labelled exploratory.

The remainder of the paper is organised as follows. Section II situates this work against open-loop conformal prediction for driving, task-relevant failure detection, and closed-loop safety evaluation. Section III formalises the triggering problem and the certification procedure. Section IV describes the closed-loop evaluation environment, the fallback manoeuvre, the harm labels and the baselines. Section V reports results,

ablations and the shift experiments. Section VI states the scope of the claims, and Section VII concludes.

II. RELATED WORK

A. Conformal prediction for autonomous driving

Conformal prediction [1] constructs prediction sets with a distribution-free, finite-sample marginal coverage guarantee under exchangeability between calibration and test data, and has been applied extensively to trajectory forecasting and to safe planning under uncertainty. Lindemann *et al.* embed conformal prediction regions directly into a model predictive controller to obtain probabilistic safety guarantees, but their guarantee collapses once the deployment distribution shifts away from the calibration distribution — a limitation Rahaman *et al.* [4] address with robust conformal prediction driven by a conditional generative model of the shift, evaluated in a synthetic multi-agent corridor. SafePath [5] integrates conformal prediction into an LLM-based path-planning pipeline, providing a guarantee that a set of candidate paths contains a safe one with a user-specified probability, and reports both open-loop (nuScenes) and closed-loop (highway-env, synthetic traffic) results. In all three cases the guarantee is on the *prediction* — a region contains the truth, or a path set contains a safe path — rather than on the eventual control decision the region is used to make, and none evaluates a fallback-triggering decision on real, replayed driving logs under reactive traffic. This paper's guarantee is instead on the decision: the probability that a minimum-risk manoeuvre fires before a harmful event, evaluated end-to-end in closed loop.

B. Online and decision-level conformal methods

Adaptive Conformal Inference (ACI) [13] updates the mis-coverage level online so that long-run coverage holds even under arbitrary shift. Conformal Decision Theory (CDT) [12] carries the same idea from prediction sets to decisions: it updates a decision parameter $\lambda_{t+1} = \lambda_t + \eta(\varepsilon - \ell_t)$ so that the long-run average risk converges to ε at rate $O(1/t)$, with no distributional assumptions. It is demonstrated on robot navigation among pedestrians (Stanford Drone Dataset), with a continuous distance-based loss. Both give *asymptotic, average* guarantees for a quantity adjusted continuously as data arrive. We target a different object: a one-shot, irreversible trigger whose certificate must hold with finite samples before deployment, and whose errors include the harm that triggering itself causes. We adopt both methods as baselines (Section IV-E), applied episode-by-episode, which is the natural online counterpart of our setting.

C. Task-relevant and closed-loop failure detection

Farid *et al.* [6] propose a run-time monitor for trajectory-prediction failures that is *task-relevant* by construction: it flags a misprediction only when the resulting planning cost is anomalous relative to the predicted cost distribution, with data-free false-positive and false-negative rate bounds. Their detector is evaluated against hand-labelled scenarios in an open-loop and semi-reactive setting, and — like the conformal-prediction

literature above — does not close the loop through an actuated fallback and re-measure the outcome. Their central insight, that a task-irrelevant misprediction should not count against a detector, is reflected in this paper’s harm labelling: a collision the ego vehicle suffers while already stopped or while being struck from behind is scored as a distinct *induced-collision* cost of the fallback itself, never as a missed intervention.

Closest to our certification step, Luo *et al.* [8] calibrate the alert threshold of a warning system with conformal prediction so that the fraction of unsafe situations that raise no alert is bounded with finite samples, and demonstrate it on driving. We adopt the same principle — certify a threshold on a risk score against missed events — and do not claim it as new. What we add is what happens *after* the alert: the fallback it triggers changes the outcome (it can prevent harm, cause an unnecessary stop, or induce a rear-end collision), so miss and over-intervention rates can only be measured by closing the loop. Our contributions are an exact way to do that for every trigger, threshold and online method on identical physics, the resulting test of whether open-loop failure-detection metrics select the right trigger, and the behaviour of the certificate under a real cross-dataset shift.

D. Risk-sensitive safety analysis

Chapman *et al.* [9] develop risk-sensitive safe sets using Conditional Value-at-Risk to assess the severity, not merely the probability, of constraint violation along a stochastic trajectory, providing the risk-sensitive vocabulary this paper borrows in motivating why a mean missed-intervention rate is the appropriate quantity to control at a chosen operating point, even though a tail-severity extension (Section VI) is left to future work.

E. Closed-loop simulation infrastructure

ScenarioNet [10] unifies heterogeneous real-world driving logs — including Argoverse 2, nuScenes, nuPlan, and Lyft Level 5 — into a common scenario description that can be replayed interactively in the MetaDrive simulator, with either logged-trajectory replay or reactive rule-based (IDM) traffic. This work is built directly on that infrastructure: every scenario reported in Section V is a real logged driving scenario, replayed headless with reactive traffic, not a synthetic or hand-crafted one.

III. PROBLEM FORMULATION AND METHOD

A. Fallback triggering as selective prediction

Consider a closed-loop rollout of a scenario X under a fallback policy governed by a scalar threshold λ : at every decision tick the vehicle computes an uncertainty score $u_t(X)$ from a trajectory predictor’s output for nearby agents, and engages the minimum-risk manoeuvre the first time $u_t(X) > \lambda$, remaining engaged for the rest of the episode. Define the closed-loop loss

$$L(X, \lambda) = \mathbb{1}[\text{harm in } X \text{ under } \lambda, \text{ MRM not first}], \quad (1)$$

the *missed-intervention* indicator, and the complementary over-conservatism quantity

$$S(X, \lambda) = \mathbb{1}[\text{MRM fires under } \lambda, X \text{ harm-free at } \lambda = \infty], \quad (2)$$

the *unnecessary stop* indicator, evaluated against the counterfactual rollout of the same scenario with the trigger disabled. The goal is to choose $\hat{\lambda}$ such that

$$\Pr_{X_1, \dots, X_n} \left(\mathbb{E}_X [L(X, \hat{\lambda})] \leq \alpha \right) \geq 1 - \delta, \quad (3)$$

while minimising $\mathbb{E}_X [S(X, \hat{\lambda})]$ — the smallest autonomy loss compatible with the risk target α .

Equation (3) is not a conformal coverage statement: $L(X, \lambda)$ depends on the entire closed-loop rollout under λ , since engaging the fallback changes every subsequent state, so it is neither a marginal coverage event on a static prediction nor guaranteed to be monotone in λ a priori. We therefore calibrate $\hat{\lambda}$ with Learn-then-Test [2]: for a finite grid $\lambda_1 > \lambda_2 > \dots > \lambda_J$ ordered from least to most conservative, compute the empirical risk $\hat{R}(\lambda_j) = \frac{1}{n} \sum_i L(X_i, \lambda_j)$ on n calibration scenarios, a Hoeffding–Bentkus p -value for the null hypothesis $H_j : \mathbb{E}[L(\cdot, \lambda_j)] > \alpha$, and apply fixed-sequence testing starting from the most conservative λ ; every λ_j whose null is rejected satisfies (3) simultaneously, with family-wise error controlled at δ , and $\hat{\lambda}$ is the least conservative such λ_j (largest autonomy). Where the empirical loss is observed to be monotone in λ across the calibration set, we additionally report Conformal Risk Control [3], whose expectation guarantee is tighter but requires that monotonicity assumption; where it is violated we report the violation rate directly rather than silently falling back to the weaker method.

B. Trigger score

At each decision tick the trigger score u_t measures how far a predicted future position of any nearby agent intrudes into the ego vehicle’s own near-term extent. Writing the ego’s constant-heading, constant-speed extrapolation over a horizon of T steps as $\{p_\tau^{\text{ego}}\}_{\tau=1}^T$ — the same route-agnostic motion model used by the harm metric in Section IV-B, deliberately *not* the vehicle’s planned, route-following path, so that the trigger cannot credit itself for steering behaviour the safety label does not assume — and a predictor’s K -mode forecast for agent i as $\{\hat{p}_\tau^{(i,k)}\}$, each predicted offset is decomposed into a longitudinal component $a_\tau^{(i,k)}$ (along the ego heading) and a lateral component $\ell_\tau^{(i,k)}$ (perpendicular to it), each compared against a separate margin: a length-based longitudinal margin $m_i^{\text{lon}} = \frac{1}{2}(L_{\text{ego}} + L_i) + c_{\text{lon}}$ and a width-based lateral margin $m_i^{\text{lat}} = \frac{1}{2}(W_{\text{ego}} + W_i) + c_{\text{lat}}$. The intrusion at (i, k, τ) is

$$\rho_\tau^{(i,k)} = \max \left(0, 1 - \sqrt{\left(\frac{a_\tau^{(i,k)}}{m_i^{\text{lon}}} \right)^2 + \left(\frac{\ell_\tau^{(i,k)}}{m_i^{\text{lat}}} \right)^2} \right) m_i^{\text{lat}}, \quad (4)$$

and $u_t = \max_{i,k,\tau} \rho_\tau^{(i,k)}$. A single isotropic clearance radius cannot represent both margins at once: a radius set by width alone risks missing a straight-ahead lead vehicle, and one set by the half-diagonal risks triggering on adjacent-lane traffic.

The decomposition in (4) is therefore a modelling choice made during development on this reasoning. Our ablation (Section V-C) finds that it gives no measurable benefit over a width-only isotropic radius, so we report it as a design option rather than a requirement.

C. Fallback manoeuvre

On triggering, the vehicle executes the minimum-risk manoeuvre defined by UN Regulation No. 157 (Automated Lane Keeping Systems) § 5.5.1 [18]: decelerate inside the current lane with a deceleration demand not exceeding 4.0 m/s^2 , lane-keeping control otherwise unchanged, holding the vehicle stationary once stopped. This is enforced by a proportional controller tracking the target speed profile $v(t) = \max(v_0 - 4.0t, 0)$ from the moment of triggering.

D. Certification under covariate shift

The certificate in (3) assumes calibration and deployment scenarios are exchangeable. When the deployment distribution differs by covariate shift, we reweight calibration scenarios by an estimated density ratio $w(x) \propto p_{\text{target}}(x)/p_{\text{source}}(x)$ computed from pre-deployment scenario covariates only (agent counts and speeds, initial gap, route length). A probabilistic domain classifier is trained on unlabelled target scenarios, following weighted conformal prediction [14]. The weighted empirical risk $\hat{R}_w(\lambda) = \sum_i w_i L(X_i, \lambda) / \sum_i w_i$ is tested with the Hoeffding–Bentkus p -value at the Kish effective sample size $n_{\text{eff}} = (\sum_i w_i)^2 / \sum_i w_i^2$. As the weights concentrate, n_{eff} shrinks and the test certifies a more conservative threshold, and eventually none at all. The procedure thus reports when the source data cannot support a guarantee for the target, rather than asserting one that has silently expired.

E. Exact counterfactual evaluation

Two rollouts of the same scenario are identical until one of them triggers the fallback, up to the simulator’s own non-determinism. We measure this in every forced-firing rollout. Over the first 33,761 forced-firing rollouts (2,000 scenarios), 98.4% agreed with the reference run to within 1 cm before firing. The pre-fire time-to-collision trace differed in 1.2% and the pre-fire harm status in one rollout (0.003%). We report the result of every analysis both on all scenarios and after dropping the 3.6% of scenarios with any flagged drift. A rollout’s outcome is therefore a function of the *firing tick* alone, not of which score, threshold or update rule caused the firing. For each scenario we simulate one reference run with the fallback disabled, recording every trigger score at every decision tick, and one rollout forced to fire at each decision tick. Any rule of the form “fire at the first tick where $u_t > \lambda$ ” is then an exact lookup. This includes online rules, whose λ changes between scenarios, and added actuation latency. All methods in Section V are compared on the same simulated physics, so differences between them cannot come from simulator noise or from unequal compute budgets.

IV. EXPERIMENTAL SETUP

A. Simulation environment

All rollouts are executed headless in MetaDrive, replaying scenarios converted by ScenarioNet [10] from Argoverse 2 Motion Forecasting and nuScenes. Non-ego traffic is controlled reactively (rule-based, IDM), not replayed open-loop from the log: pure log replay was found to roughly triple the no-fallback collision rate on a 1,000-scenario diagnostic sample, because logged agents cannot react to an ego vehicle that deviates from its logged trajectory — precisely the deviation a fallback manoeuvre produces. Reporting results under non-reactive replay would therefore misrepresent the cost of stopping, which is a central quantity this paper reports.

B. Harm labelling

A rollout is labelled harmful using the nuPlan closed-loop safety definition [11]: an at-fault collision, or a time-to-collision below 0.95 s computed over a 3.0 s horizon at 0.1 s resolution, not evaluated while the ego is stationary and not evaluated against tracks behind the ego. Fault attribution follows the nuPlan convention: a collision the ego suffers while stationary, or while being struck from behind, is *not* at-fault and is never counted as a missed intervention; it is instead recorded as an *induced collision* — the concrete, measurable cost of a false or premature trigger.

C. Predictor

The trajectory predictor is AutoBot [16] (6 modes, 2.1 s history, 6 s horizon), trained with the UniTraj framework [17] on the Argoverse 2 Motion Forecasting training split (minADE₆ = 1.09 m on held-out AV2 validation scenes). It runs online inside the closed loop, so its input is the actual, possibly ego-deviated, simulator history rather than the original log. This is necessary once traffic is reactive. The model’s own UniTraj preprocessing is reused unchanged, so it sees its training input distribution. Vehicles are predicted by the network. Pedestrians and cyclists, which the checkpoint was not trained on, receive constant-velocity forecasts.

D. Datasets and splits

All scenarios come from Argoverse 2 validation, disjoint from the predictor’s training data. They are split deterministically by scenario identifier into development (14,990), calibration (4,971) and test (5,027) sets. The shifted target is nuScenes (9,041 scenarios of 8.1 s, Boston and Singapore, versus Argoverse 2’s six U.S. cities). Scenarios in which the ego moves less than 10 m are excluded, as in ScenarioNet. Design choices were made on the development split only. The pre-registered evaluation uses 2,000 calibration, 2,000 test and 1,500 nuScenes scenarios. Ten Argoverse 2 scenarios (six calibration, four test) whose per-agent records could not be matched after a re-scoring pass are dropped from every analysis, so all methods are compared on identical scenarios; this leaves 1,994 calibration and 1,996 test scenarios, and 1,504 nuScenes scenarios. A second shifted target, Waymo Open Motion [7] validation scenarios (9.1 s, 1 s of history), is

collected under a separate pre-registered amendment [results pending]. Decisions are taken at 2 Hz. A paired development-split comparison showed identical miss rates at 2 Hz and 10 Hz, and a 10 Hz replication on 100 test scenarios is reported.

E. Baselines

All methods use the same predictor, the same MRM and the same test scenarios, and are evaluated exactly (Section III-E).

- **T1 — confidence threshold** (deployed practice): fire when $1 - \max_k p_k$ of the network-predicted vehicle whose predicted modes come closest to the ego exceeds a threshold. Agents the network does not predict (pedestrians, cyclists, distant vehicles) are excluded from this score. The threshold is tuned on calibration data to the most permissive value with empirical miss rate $\leq \alpha$. A variant that uses the nearest agent ahead is also evaluated.
- **T2 — ensemble threshold**: the same tuning, applied to the spread of endpoint predictions across five Monte Carlo dropout passes [15].
- **T3 — open-loop conformal**: inflate every predicted mode by the split-conformal radius of the predictor's 3 s error on held-out Argoverse 2 logs at level α [1], and fire when the inflated region touches the ego. This is the guarantee-on-the-prediction design that the paper argues is insufficient.
- **Tuned geometric**: our score with an empirically tuned threshold and no certificate. It isolates the contribution of the calibrator.
- **CDT** [12] and **ACI** [13]: online decision- and prediction-level adaptation, updated after every scenario. Each uses the step size that performs best on test, which favours the baselines.
- **Oracle**: the best threshold in hindsight on test data. It is not deployable and is reported as an upper bound.

F. Metrics and statistics

For every method we report the missed-intervention, unnecessary-stop and fallback-induced-collision rates, autonomy (the fraction of scenarios completed without the fallback) and route completion, with Wilson 95 % intervals. Validity of each certificate is assessed over 200 random resplits of the pooled calibration and test scenarios. A valid certificate keeps the frequency of test miss rate above α at or below δ . Harm is scored with the single nuPlan definition above; its sensitivity to the time-to-collision threshold is not evaluated here (Section VI).

G. Pre-registration and reproducibility

The harm definition, $\alpha = 0.05$, $\delta = 0.10$, and the refutation criteria for three hypotheses were recorded in a version-controlled note before any calibration or test data were collected: **H1**, that open-loop metrics rank triggers differently from closed-loop outcomes; **H2**, that a certified trigger achieves a better safety-conservatism trade-off than a tuned fixed threshold; and **H3**, that a threshold certified

TABLE I
CLOSED-LOOP OUTCOMES ON THE ARGOVERSE 2 TEST SET AT $\alpha = 0.05$.
RATES IN %. LAST COLUMN: FRACTION OF 200 CALIBRATION/TEST
RESPLITS IN WHICH THE TEST MISS RATE EXCEEDS α (VALID IF ≤ 10).

Method	Miss	Unnec.	Induced	Auton.	Viol.
No fallback	15.1	0.0	0.0	100	—
Always fire	0.0	84.9	6.5	0	—
T1 confidence	4.6	68.3	4.0	20.3	50.5
T2 MC-dropout spread	4.0	70.8	4.9	17.6	—
T3 open-loop conformal	0.0	73.8	5.5	11.1	0.0
Tuned geometric (no cert.)	4.3	26.0	2.3	61.4	43.5
CRC (geometric)	4.3	26.0	2.3	61.4	41.0
LTT (geometric, certified)	3.7	28.5	2.6	58.5	9.5
CDT (online)	4.9	24.1	2.2	63.9	—
ACI (online)	1.3	45.0	3.5	40.2	—
Oracle (hindsight)	5.0	23.6	2.2	64.2	—

on Argoverse 2 loses its guarantee on nuScenes. The note carries dated amendments (a corrected confidence baseline, the exactness measurements, a Waymo target and a reverse-shift analysis), each written before the corresponding outcomes were evaluated, and it labels the analyses run afterwards as post hoc. Code, per-scenario rollout records and configuration files will be released with the paper. Each scenario's record contains every score trace and outcome, so every table can be regenerated without re-simulation.

V. RESULTS

All results use $\alpha = 0.05$ and $\delta = 0.10$, 1,994 calibration and 1,996 test scenarios from Argoverse 2, and 1,504 nuScenes scenarios, with Wilson 95 % intervals unless stated. Dropping the 3.6 % of scenarios with any flagged simulator drift changes no conclusion below. Sections V-A–V-D are the pre-registered analyses. Section V-E combines a pre-registered result with a post hoc sensitivity analysis, and Section V-F is entirely post hoc; both are labelled.

A. Does a certificate keep its promise, and what does it cost? (RQ1)

Table I gives every method at its operating point on the Argoverse 2 test set. The certified trigger (LTT) misses 3.7 % of harmful events [95 % CI 2.9–4.6], below the 5 % target, with 28.5 % unnecessary stops and 2.6 % fallback-induced collisions. The last column repeats the calibration on 200 random resplits and reports how often the realised test miss rate exceeds α . A valid certificate keeps this at or below $\delta = 10$ %. LTT does (9.5 %). The same geometric score with a threshold tuned empirically, and its Conformal Risk Control (CRC) counterpart, which guarantees only an expectation, exceed the target in 43.5 % and 41.0 % of resplits, and the tuned confidence trigger in 50.5 %. Certification is therefore the only design here that keeps its target reliably.

The certificate has a price, and it does not win the trade-off. The tuned geometric threshold stops less often than LTT (26.0 % versus 28.5 %; paired bootstrap difference -2.5 points, 95 % CI $[-3.2, -1.8]$), and CDT stops less still (24.1 %). Our pre-registered hypothesis H2 required no baseline to be within 3 points of the certified trigger; it is **refuted**. The supportable claim is narrower: among the designs

TABLE II
OPEN-LOOP AUROC VERSUS CLOSED-LOOP UNNECESSARY STOPS AT
THE CERTIFIED THRESHOLD, ARGOVERSE 2 TEST.

Score	AUROC	Unnec. stops (%)
Geometric, isotropic radius	0.843	27.5
Geometric, length-aware box	0.832	28.5
Predictor gap	0.817	31.8
Geometric, route corridor	0.577	84.9
Confidence	0.448	75.8
Ensemble spread	0.442	72.3

whose promise can be checked, the certified one is the only one that reliably meets it, and it pays about 2.5 points of unnecessary stops for that.

Deployed-practice triggers are far more conservative. Tuned to the same miss target, the confidence and MC-dropout triggers stop in 68.3 % and 70.8 % of scenarios, 40–45 points more than the certified geometric trigger (paired CIs exclude zero by a wide margin), and the open-loop conformal trigger reaches a miss rate of 0.0 % only by firing in 73.8 %. Its certificate is valid in the sense that it never violates, and useless in the sense that a vehicle that stops in three of four scenarios has abandoned its task. ACI reaches 1.3 % misses at 45 % stops, which is consistent with a guarantee on the prediction not translating into a useful decision.

B. Do open-loop metrics choose the closed-loop trigger? (RQ2)

Table II ranks six trigger scores by open-loop AUROC of their scenario peak for predicting harm, and by closed-loop unnecessary stops at each score’s certified threshold. Rankings agree at the top and diverge below it: Kendall’s $\tau = 0.60$ on Argoverse 2 and 0.47 on nuScenes, and the top-ranked score is the same on Argoverse 2 and differs on nuScenes only by 0.2 points of stops. The pre-registered rule refutes H1 only if $\tau \geq 0.8$ and the top score coincides in every setting, so H1 is not refuted, but the evidence is weak: open-loop AUROC picked the closed-loop best score, or a statistical tie, in every setting. The clearest disagreement is at the bottom. The confidence score has AUROC 0.45, below chance, meaning higher confidence-based uncertainty is on average associated with *less* harm, yet a threshold on it can still reach the miss target, at 75.8 % stops (certified) rather than by discriminating harmful from safe scenarios.

C. Ablations and latency (RQ4)

Table III shows that two components matter and one does not. Using the vehicle’s planned, route-following path instead of a kinematic corridor makes the trigger fire in essentially every scenario (84.9 % stops), consistent with the route reaching conflict points that the harm label, which uses a kinematic extrapolation, does not count. Replacing geometry with confidence or ensemble spread costs 44–47 points of stops. The length-aware box does not matter: a width-only isotropic radius performs the same (27.5 % versus 28.5 %). Delaying the intervention by one decision tick (0.5 s) already lifts the miss rate above the target (5.3 %) and two ticks give

TABLE III
ABLATIONS ON THE ARGOVERSE 2 TEST SET, CERTIFIED BY LTT.

Variant	Miss (%)	Unnec. stops (%)
Full score, LTT, no latency	3.7	28.5
Route-following ego corridor	0.0	84.9
Width-only isotropic radius	3.6	27.5
Confidence instead of geometry	3.6	75.8
Ensemble spread instead of geometry	3.6	72.3
CRC instead of LTT	4.3	26.0
Empirical tuning instead of LTT	4.3	26.0
+1 decision tick latency (0.5 s)	5.3	26.3
+2 decision ticks latency (1.0 s)	9.1	23.4

TABLE IV
A THRESHOLD CERTIFIED ON ARGOVERSE 2 (LTT, $\alpha = 0.05$) DEPLOYED
ON nuSCENES. “AVOIDABLE” EXCLUDES THE 5.0 % OF nuSCENES
SCENARIOS WITH HARM AT THE FIRST DECISION TICK (POST HOC).

Setting	Miss (%) [95 % CI]	Unnec. (%)	n
Argoverse 2 test (in domain)	3.7 [2.9, 4.6]	28.5	1,996
nuScenes, unweighted	13.2 [11.6, 15.0]	26.1	1,504
nuScenes, weighted (no labels)	7.0 [5.8, 8.4]	52.4	1,504
nuScenes, oracle threshold	5.0	77.9	1,504
Avoidable, unweighted	9.2 [7.8, 10.8]	27.1	1,429
Avoidable, weighted ($n_{\text{eff}} = 68$)	2.2 [1.6, 3.1]	54.5	1,429

9.1 %, while stops fall, so a certificate calibrated at zero latency is not valid for a slower actuation chain and latency belongs in the certified quantity.

D. Dataset shift (RQ3)

The Argoverse 2 certificate does not transfer. On nuScenes the threshold certified at a 5 % target misses 13.2 % [11.6, 15.0], 2.6 times the target and with an interval that excludes it; the pre-registered hypothesis H3 is **supported** (15.2 % after dropping drift scenarios). Part of this is a floor rather than shift: 5.0 % of nuScenes scenarios already contain harm at the first decision tick, which even a trigger that always fires cannot prevent (0.1 % on Argoverse 2). That floor equals the target, so no threshold can be certified at $\alpha = 0.05$ on nuScenes. With those scenarios excluded (post hoc), the shifted miss rate is 9.2 % [7.8, 10.8], still above the target.

Density-ratio weighting, using only unlabelled target covariates, moves the miss rate from 13.2 % to 7.0 %, still above the target, and doubles unnecessary stops (26.1 % to 52.4 %). On the avoidable subset it reaches 2.2 %, below the target, but its effective sample size is only 68, so the certificate is achieved almost entirely through conservatism (54.5 % stops).

E. Reverse shift (exploratory, post hoc)

We also certified on half of the nuScenes scenarios and deployed on Argoverse 2 test, a direction pre-registered as a reverse-shift analysis. As pre-registered, with all scenarios, the procedure *refused to certify* on both halves, because of the 5.0 % floor. On the avoidable subset (post hoc), thresholds certified on nuScenes ($n \approx 715$) give an Argoverse 2 miss rate of 1.9 % and 2.0 %, both below the target, at 40.6 % and 39.5 % unnecessary stops. A certificate from the harder dataset therefore transfers safely but conservatively, while the reverse

TABLE V
EXPLORATORY: RECALIBRATION WITH k LABELLED nuSCENES SCENES
(AVOIDABLE SUBSET, 60 RANDOM DRAWS). VIOL.: FRACTION OF DRAWS
WITH TEST MISS RATE ABOVE 5 %.

k	Method	Miss (%)	Viol. (%)	Unnec. (%)
0	AV2-certified	9.2	–	27.1
40	Target-only tuned	6.6	67	35.6
40	Target-only LTT	0.0	0	81.3
40	AV2 + k weighted LTT	1.8	10	63.6
60	Target-only tuned	6.0	70	36.7
60	Target-only LTT	0.8	0	72.5
60	AV2 + k weighted LTT	2.8	0	53.3

does not transfer. This asymmetry is a single-pair observation pending replication.

F. Recalibration with a few labelled target scenes (exploratory)

Table V draws k labelled nuScenes scenes at random (60 draws each) and evaluates on the rest, on the avoidable subset. This analysis was not pre-registered. Uncertified tuning on 40–60 target scenes exceeds the target in about two draws in three. A certified threshold on the same scenes keeps the target (violations 0–10 %) but stops in 53–81 % of scenarios, against 27 % without target labels. A handful of labelled scenes therefore restores a *valid* certificate, not the original operating point.

G. Second shifted dataset and sensitivity

Waymo Open Motion. [Pending: miss rate of the Argoverse 2-certified threshold on Waymo (pre-registered as H3-W: the 95 % interval excludes α), weighted repair, and open-loop versus closed-loop ranking.]

Sensitivity. [Pending: non-reactive log-replay traffic, a 2.0 m/s² comfort MRM, and 10 Hz decisions on 100 scenarios.]

VI. SCOPE AND FUTURE WORK

The claims above are deliberately scoped. Each boundary below is a design choice with a stated rationale, and each suggests a direct extension.

What we claim, and what we do not. We claim that certification is what makes a closed-loop trigger keep its stated miss rate, that it costs a measurable number of unnecessary stops, and that the guarantee is fragile to dataset shift and to latency. We do not claim a better-performing trigger: a tuned threshold and an online method stop less often, and we report the pre-registered hypothesis that said otherwise as refuted.

One predictor, one simulator. All results use one trajectory predictor trained on Argoverse 2, and MetaDrive with rule-based (IDM) agents. Whether the ordering of triggers survives a stronger predictor or learned agent behaviour is open. The protocol accepts both unchanged, and the geometric-versus-confidence gap is large enough that we expect it to persist, which we have not tested.

Unavoidable harm. The 5.0 % first-tick floor on nuScenes shows that a miss metric defined on harm anywhere in a

replayed snippet mixes events a fallback could prevent with events already under way. We report the pre-registered definition and the avoidable-only variant side by side, and propose the latter as a refinement for future benchmarks.

Marginal, scenario-level guarantees. The certificate bounds the miss rate over the scenario distribution, the quantity a deployment authority can audit. It does not bound the risk of any individual scenario. Stratified certification by scenario type follows from the same procedure applied within strata.

Covariate shift and few-shot recalibration. The weighted certificate covers shift in observable scenario covariates, and its effective sample size, 68 in our nuScenes subset, shows when it becomes unreliable. The few-label recalibration result is exploratory and single-pair; it needs replication on a second target before being relied on.

Harm definition. Harm is scored with the nuPlan definition alone. Sensitivity to the time-to-collision threshold and to fault attribution is a natural robustness study, and a severity-weighted or tail (CVaR) certificate in the sense of [9] requires only a change of loss.

Object states as input; one fallback family. The predictor receives agent states, isolating the triggering decision from perception error, and we evaluate stop-in-lane manoeuvres at the UN R157 deceleration limit. Composing the certificate with a calibrated perception module, and other fallbacks such as a lateral pull-over or hand-over, change the quantities priced but not the method.

VII. CONCLUSION

A safety guarantee on a fallback trigger is not a permanent property of the model; it depends on the data it was certified on, the timing of the intervention, and the consequences of the fallback itself. Evaluated exactly in closed loop, a Learn-then-Test-certified trigger keeps its 5 % target on the calibration distribution (3.7 % missed, violations in 9.5 % of resplits against 43.5 % without a certificate) at a cost of about 2.5 points of unnecessary stops, and deployed-practice confidence and ensemble triggers need 68–71 % stops for similar safety. The same certificate fails on nuScenes (13.2 % missed), is broken by one decision tick of latency, and can be unreachable when the target data contain harm no trigger can prevent. Recalibrating with a few dozen labelled target scenes restores a valid certificate at a large operating cost in an exploratory analysis. [Waymo replication to be added.] These findings argue for treating certification, shift monitoring and latency as one audited quantity, and the exact evaluation protocol makes that audit inexpensive to repeat for any trigger.

ACKNOWLEDGMENT

The authors thank the maintainers of MetaDrive and ScenarioNet for releasing the simulation infrastructure this work builds on.

REFERENCES

- [1] V. Vovk, A. Gammernan, and G. Shafer, *Algorithmic Learning in a Random World*. Berlin, Germany: Springer-Verlag, 2005.

- [2] A. N. Angelopoulos, S. Bates, E. J. Candès, M. I. Jordan, and L. Lei, “Learn then test: Calibrating predictive algorithms to achieve risk control,” *arXiv preprint arXiv:2110.01052*, 2021.
- [3] A. N. Angelopoulos, S. Bates, A. Fisch, L. Lei, and T. Schuster, “Conformal risk control,” *arXiv preprint arXiv:2208.02814*, 2022.
- [4] K. Rahaman, J. V. Deshmukh, A. R. Hota, and L. Lindemann, “When environments shift: Safe planning with generative priors and robust conformal prediction,” in *Proc. 8th Annu. Conf. Learn. Dyn. Control (LADC)*, 2026.
- [5] A. Doula, M. Mühlhäuser, and A. S. Guinea, “SafePath: Conformal prediction for safe LLM-based autonomous navigation,” *arXiv preprint arXiv:2505.09427*, 2025.
- [6] A. Farid, S. Veer, B. Ivanovic, K. Leung, and M. Pavone, “Task-relevant failure detection for trajectory predictors in autonomous vehicles,” in *Proc. 6th Conf. Robot Learn. (CoRL)*, Auckland, New Zealand, 2022.
- [7] S. Ettinger *et al.*, “Large scale interactive motion forecasting for autonomous driving: The Waymo Open Motion Dataset,” in *Proc. IEEE/CVF Int. Conf. Comput. Vis. (ICCV)*, 2021, pp. 9710–9719.
- [8] R. Luo, S. Zhao, J. Kuck, B. Ivanovic, S. Savarese, E. Schmerling, and M. Pavone, “Sample-efficient safety assurances using conformal prediction,” in *Proc. Workshop Algorithmic Found. Robot. (WAFR)*, 2022.
- [9] M. P. Chapman, R. Bonalli, K. M. Smith, I. Yang, M. Pavone, and C. J. Tomlin, “Risk-sensitive safety analysis using conditional value-at-risk,” *IEEE Trans. Autom. Control*, vol. 67, no. 12, pp. 6521–6536, Dec. 2022.
- [10] Q. Li, Z. Peng, L. Feng, Z. Liu, C. Duan, W. Mo, and B. Zhou, “ScenarioNet: Open-source platform for large-scale traffic scenario simulation and modeling,” in *Proc. 37th Conf. Neural Inf. Process. Syst. (NeurIPS), Track Datasets and Benchmarks*, 2023.
- [11] N. Karnchanachari *et al.*, “Towards learning-based planning: The nuPlan benchmark for real-world autonomous driving,” in *Proc. IEEE Int. Conf. Robot. Autom. (ICRA)*, 2024.
- [12] J. Lekeufack, A. N. Angelopoulos, A. Bajcsy, M. I. Jordan, and J. Malik, “Conformal decision theory: Safe autonomous decisions from imperfect predictions,” in *Proc. IEEE Int. Conf. Robot. Autom. (ICRA)*, 2024, pp. 11668–11675.
- [13] I. Gibbs and E. Candès, “Adaptive conformal inference under distribution shift,” in *Proc. Adv. Neural Inf. Process. Syst. (NeurIPS)*, vol. 34, 2021, pp. 1660–1672.
- [14] R. J. Tibshirani, R. Foygel Barber, E. J. Candès, and A. Ramdas, “Conformal prediction under covariate shift,” in *Proc. Adv. Neural Inf. Process. Syst. (NeurIPS)*, vol. 32, 2019.
- [15] Y. Gal and Z. Ghahramani, “Dropout as a Bayesian approximation: Representing model uncertainty in deep learning,” in *Proc. 33rd Int. Conf. Mach. Learn. (ICML)*, 2016, pp. 1050–1059.
- [16] R. Girgis, F. Golemo, F. Codevilla, M. Weiss, J. A. D’Souza, S. E. Kahou, F. Heide, and C. Pal, “Latent variable sequential set transformers for joint multi-agent motion prediction,” in *Proc. Int. Conf. Learn. Represent. (ICLR)*, 2022.
- [17] L. Feng, M. Bahari, K. M. Ben Amor, É. Zablocki, M. Cord, and A. Alahi, “UniTraj: A unified framework for scalable vehicle trajectory prediction,” in *Proc. Eur. Conf. Comput. Vis. (ECCV)*, 2024.
- [18] United Nations Economic Commission for Europe, “UN Regulation No. 157 — Uniform provisions concerning the approval of vehicles with regard to Automated Lane Keeping Systems,” Rev. 1, 2025.